

Hotel Reservation

By Hantao Lin





Problem

According <http://blog.experience-hotel.com/>, the average rate of **hotel cancellation from all sources is about 24%**. Consequently, it has always been the hotel revenue manager's biggest concern to counteract hotel cancelation. The mission here is to **build an appropriate model that predicts the likelihood of hotel cancelation** that helps revenue managers to better understand the situations and propose policies to accommodate or avoid hotel cancelation.



Context of the Data

There are two datasets used in this project:

1. Kaggle Dataset (<https://www.kaggle.com/jessemostipak/hotel-booking-demand>)
 - The dataset contains booking information for a city hotel and a resort hotel.
 - Includes information such as when the booking was made, length of stay, the number of adults, children, and/or babies, and the number of available parking spaces, among other 32 features.
 - It has 32 columns and 119390 entries.

```

print(hotel.info())

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 119390 entries, 0 to 119389
Data columns (total 32 columns):
hotel                                119390 non-null object
is_canceled                          119390 non-null int64
lead_time                            119390 non-null int64
arrival_date_year                     119390 non-null int64
arrival_date_month                    119390 non-null object
arrival_date_week_number              119390 non-null int64
arrival_date_day_of_month             119390 non-null int64
stays_in_weekend_nights              119390 non-null int64
stays_in_week_nights                 119390 non-null int64
adults                               119390 non-null int64
children                             119386 non-null float64
babies                               119390 non-null int64
meal                                  119390 non-null object
country                              118902 non-null object
market_segment                       119390 non-null object
distribution_channel                  119390 non-null object
is_repeated_guest                    119390 non-null int64
previous_cancellations                119390 non-null int64
previous_bookings_not_canceled        119390 non-null int64
reserved_room_type                    119390 non-null object
assigned_room_type                    119390 non-null object
booking_changes                       119390 non-null int64
deposit_type                          119390 non-null object
agent                                103050 non-null float64
company                              6797 non-null float64
days_in_waiting_list                 119390 non-null int64
customer_type                         119390 non-null object
adr                                   119390 non-null float64
required_car_parking_spaces           119390 non-null int64
total_of_special_requests              119390 non-null int64
reservation_status                    119390 non-null object
reservation_status_date               119390 non-null object
dtypes: float64(4), int64(16), object(12)
memory usage: 29.1+ MB
None

```

Figure 1. Kaggle Data



2. Wunderground data (wunderground.com)

- The dataset includes the average temperature, wind speed, precipitation and other things at Lisbon,
- It is collected from 7/1/2015 to 8/31/2017.
- It has 9 columns and 793 entries.

```
print(weather.info())
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 793 entries, 0 to 792  
Data columns (total 9 columns):  
Date_time                793 non-null object  
Maximum_Temperature      793 non-null float64  
Minimum_Temperature      793 non-null float64  
Average_Temperature      793 non-null float64  
Heat_Index               184 non-null float64  
Precipitation(inches)    793 non-null float64  
Wind_Speed               793 non-null float64  
Relative_Humidity        793 non-null float64  
Conditions               793 non-null object  
dtypes: float64(7), object(2)  
memory usage: 55.9+ KB  
None
```

Figure 2. Weather Data



Data Wrangling

The datasets were joined for the analysis. In order to do so, a new column(“Date_time”) is created in Kaggle dataset based on the “arrival_year”, “arrival_month”, and “arrival_day” in the dataset. After joining the two datasets with “Date_time” as the common column, missing values were filled differently in the new dataset. Missing value in “Heat_Index” is calculated based on a formula provided online. Other missing values in other columns are filled with the mode of the column. Column “company” is dropped from the dataset because it is missing 80% of its values. After the cleaning, the final dataset has 40 columns in total and 119390 entries.

```

merged.info()

<class 'pandas.core.frame.DataFrame'>
Int64Index: 119390 entries, 0 to 119389
Data columns (total 41 columns):
hotel                119390 non-null object
is_canceled          119390 non-null int64
lead_time            119390 non-null int64
arrival_date_year     119390 non-null int64
arrival_date_month    119390 non-null object
arrival_date_week_number 119390 non-null int64
arrival_date_day_of_month 119390 non-null int64
stays_in_weekend_nights 119390 non-null int64
stays_in_week_nights  119390 non-null int64
adults               119390 non-null int64
children             119390 non-null float64
babies               119390 non-null int64
meal                 119390 non-null object
country              119390 non-null object
market_segment       119390 non-null object
distribution_channel  119390 non-null object
is_repeated_guest    119390 non-null int64
previous_cancellations 119390 non-null int64
previous_bookings_not_canceled 119390 non-null int64
reserved_room_type    119390 non-null object
assigned_room_type    119390 non-null object
booking_changes       119390 non-null int64
deposit_type          119390 non-null object
agent                119390 non-null float64
days_in_waiting_list 119390 non-null int64
customer_type         119390 non-null object
adr                  119390 non-null float64
required_car_parking_spaces 119390 non-null int64
total_of_special_requests 119390 non-null int64
reservation_status    119390 non-null object
reservation_status_date 119390 non-null object
Date_time             119390 non-null object
Maximum_Temperature   119390 non-null float64
Minimum_Temperature    119390 non-null float64
Average_Temperature    119390 non-null float64
Heat_Index             119390 non-null float64
Precipitation(inches)  119390 non-null float64
Wind_Speed             119390 non-null float64
Relative_Humidity      119390 non-null float64
Conditions             119390 non-null object
Heat_index             119390 non-null float64
dtypes: float64(11), int64(16), object(14)
memory usage: 43.3+ MB

```

Figure 3. Merged Data



Data Storytelling

After getting the cleaned data, I had some initial guesses about the factors that influence the chance of people canceling their hotel appointment :

1. Whether weather such as average temperature and wind speed have effects on the cancelation rate?
2. How is status of previous appointments correlated with cancelation rate?
3. If there are countries that have an extremely higher cancelation rate compared to the others?

Cancellation Rate vs Average Temperature

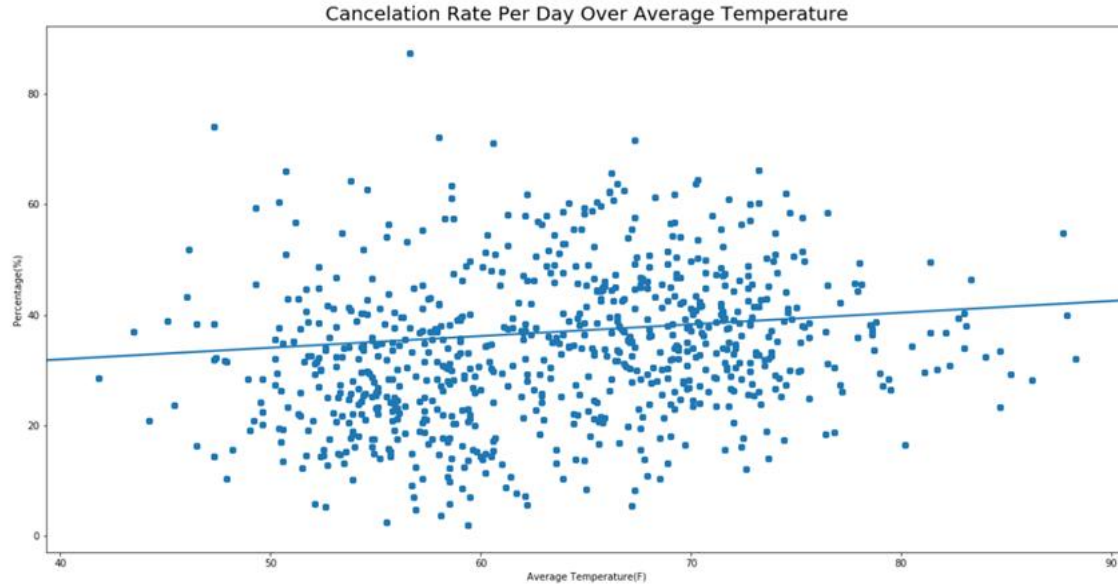


Figure 4. Cancellation Rate Per Day Over Average Temperature

From the graph, we can see there is a positive correlation between average temperature and cancellation rate over time

Cancellation Rate vs Wind Speed





Figure 5. Cancellation Rate Per Day Over Wind Speed

From the graph above, wind speed and cancellation rate has a flat to negative regression line. This tells us that there is not much correlation between wind speed and rate of cancellation.

Cancellation Rate VS Previous Appointment Status



From the graph we can conclude that people with 14 or more times of cancellation records have a 100% chance of cancelling again. Moreover, people who canceled 1 time and 13 times before are also extremely likely to cancel their reservation.

Figure 6. Cancellation Rate Over Previous Cancellations



Cancelation Rate VS Countries

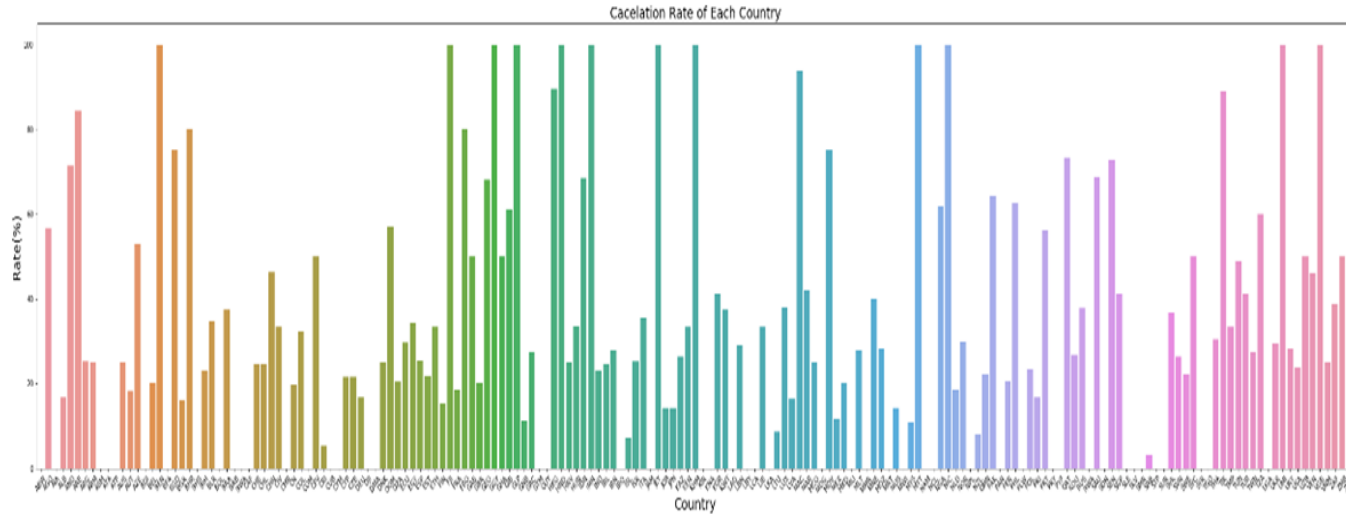


Figure 7. Cancelation Rate Over Country.
There are indeed some countries that have higher cancellation rates compared to other countries.



Statistical Analysis

In this step, I will further investigate the correlation between wind speed, average temperature, and cancelation rate. Moreover, I will expand on the graph of cancelation rate of each country by examining if the domestic (Portugal) cancelation rate is equal to the international cancelation rate



Hypothesis Test for Average Temperature VS Cancellation Rate

Null Hypothesis: Pearson's coefficient = 0

Alternative Hypothesis: Pearson's coefficient $\neq 0$

```
cor_list=[]
for i in range(10000):
    sample1=np.random.choice(d['percent'],len(d['percent']))
    sample2=np.random.choice(d['temperature'],len(d['temperature']))
    cor=stats.pearsonr(sample1,sample2)
    cor_list.append(cor)
```

```
correlation=[]
for i in cor_list:
    correlation.append(i[0])
```

```
print('Correlation of Average Temperature and Cancellation Rate 95% Confidence Interval:'
      +str(np.percentile(correlation,[2.5,97.5])))
```

Correlation of Average Temperature and Cancellation Rate 95% Confidence Interval:[-0.07026314 0.07109968]

Figure 8. Confidence Interval of Correlation of Average Temperature and Cancellations Rate

Because the 95% confidence interval contains 0, it is inconclusive whether the two

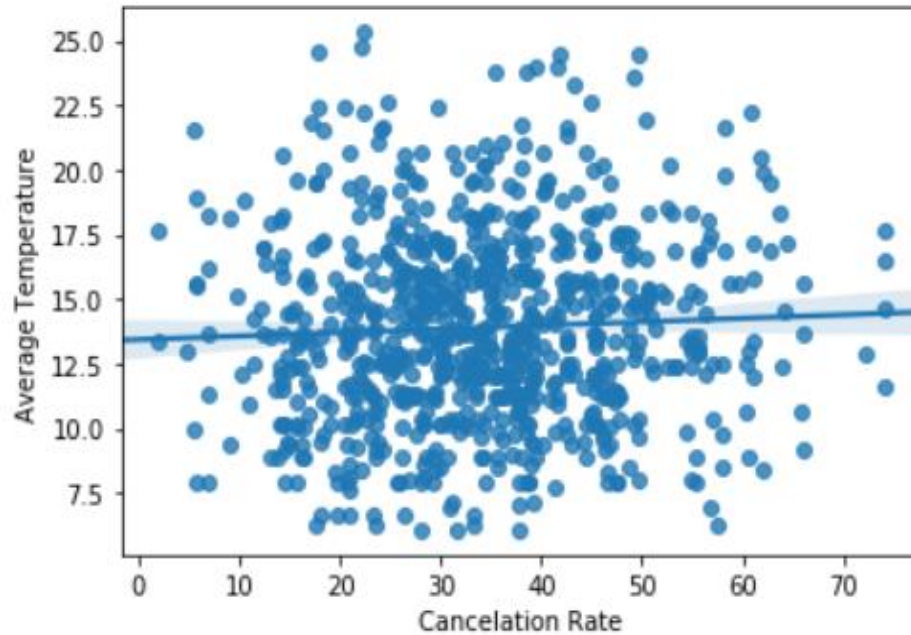


Figure 9. Bootstraps Sample of Cancellation Rate Over Average Temperature

The graph is the bootstrap points representing cancellation rate vs average temperature. It is very hard to conclude that there exists a significant correlation between the two variables since the regression line is not very steep.



Hypothesis Test for Wind Speed VS Cancellation Rate

Null Hypothesis: Pearson's coefficient = 0

Alternative Hypothesis: Pearson's coefficient \neq 0

```
cor_list2=[]
for i in range(10000):
    sample1=np.random.choice(d['percent'],len(d['percent']))
    sample2=np.random.choice(d['wind_speed'],len(d['wind_speed']))
    cor=stats.pearsonr(sample1,sample2)
    cor_list2.append(cor)
```

```
correlation2=[]
for i in cor_list2:
    correlation2.append(i[0])
```

```
p=[]
for i in cor_list2:
    p.append(i[1])
print('Correlation of Wind Speed and Cancellation Rate 95% Confidence Interval:'+
      str(np.percentile(correlation2,[2.5,97.5])))
print('p-value: '+str(np.mean(p)))
```

Correlation of Wind Speed and Cancellation Rate 95% Confidence Interval:[-0.06967139 0.06925268]
p-value:0.5026847741192151

Figure 10. Confidence Interval of Cancellation Rate and Wind Speed

The result made the test inconclusive because the confidence interval contains both negative and positive numbers. Similar to the average temperature vs cancelation rate, the correlation between wind speed and cancelation rate could be positive, negative, or not correlated at all.




Figure 11. Bootstraps Sample of Cancellation Rate Over Wind Speed

The graph is very similar to temperature vs cancellation rate. This graph also illustrates that there is not a significant correlation between wind speed and cancellation rate.



Hypothesis Test for Domestic Cancellation Rate VS International Cancellation Rate

Null Hypothesis: mean domestic cancellation rate - mean international cancellation rate = 0

Alternative Hypothesis: Pearson's coefficient: mean domestic cancellation rate - mean international cancellation rate \neq 0

```
list=[]
for i in range(10000):
    sample1=np.random.choice(dom_count['rate'],len(dom_count['rate']))
    sample2=np.random.choice(int_count['rate'],len(int_count['rate']))
    mean1=np.nanmean(sample1)
    mean2=np.nanmean(sample2)
    diff=mean1-mean2
    list.append(diff)
percentile=np.percentile(list,[2.5,97.5])
print('Domestic vs International Cancellation Rate 95% Confidence Interval:'+str(percentile))
```

Domestic vs International Cancellation Rate 95% Confidence Interval:[-0.02214565 0.02167502]

The 95% confidence interval is $(-0.02134154, 0.02174621)$ which includes 0. This made the hypothesis test inconclusive.




Figure 13. Distribution of Domestic and International Cancellation Rate

The distribution plot vividly demonstrates that it is not possible to conclude as the two distributions overlap each other heavily.